

Detailed Explanation of LSTM + DQN (and PPO) Trading Workflow

Your Name

February 24, 2025

Contents

1	Introduction	2
2	Full Source Code	2
3	Step-by-Step Explanation	9
3.1	Imports and Configuration	9
3.2	Data Loading and Preprocessing (Functions in Step 1)	9
3.3	Argument Parsing (Step 2)	10
3.4	Main Function (Step 3)	10
3.5	Building and Tuning the LSTM (Steps 7–9)	10
3.6	Evaluation and Saving the Model (Steps 10–11)	11
3.7	Reinforcement Learning Environment and DQN Training (Step 12)	11
4	Replacing DQN with PPO	11
4.1	Potential Benefits of PPO over DQN	11
4.2	Modifications in the Code	12
5	Conclusion	12

1 Introduction

This document provides a step-by-step explanation of the Python program that trains:

1. An LSTM neural network (for price prediction) using Optuna for hyperparameter tuning.
2. A Deep Q-Network (DQN) agent (from Stable-Baselines3) in a custom Stock Trading Gym environment for reinforcement learning.

Additionally, it will cover:

- The benefits of potentially substituting the DQN algorithm with Proximal Policy Optimization (PPO).
- A detailed guide on how to modify the existing code to implement a PPO agent instead of a DQN agent.

2 Full Source Code

In this section, we display the full source code. Each component will be described in detail in subsequent sections.

Listing 1: Full program source code.

```
1 import os
2 import sys
3 import argparse
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 import logging
9 from tabulate import tabulate
10
11 from sklearn.preprocessing import MinMaxScaler
12 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
13
14 import tensorflow as tf
15 from tensorflow.keras.models import Sequential
16 from tensorflow.keras.layers import LSTM, Dense, Dropout, Bidirectional
17 from tensorflow.keras.optimizers import Adam, Nadam
18 from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
19 from tensorflow.keras.losses import Huber
20 from tensorflow.keras.regularizers import l2
21
22 import xgboost as xgb
23 import optuna
24 from optuna.integration import KerasPruningCallback
25
26 # For Reinforcement Learning
27 import gym
28 from gym import spaces
29 from stable_baselines3 import DQN
30 from stable_baselines3.common.vec_env import DummyVecEnv
31
32 # Suppress TensorFlow warnings
33 os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2' # Suppress INFO/WARNING
34
35 # Configure logging
36 logging.basicConfig(level=logging.INFO, format='%(asctime)s - %(levelname)s - %(message)s')
37
38 #####
39 # 1. Data Loading / Indicators
40 #####
41
42 def load_data(file_path):
43     logging.info(f"Loading data from: {file_path}")
44     try:
45         df = pd.read_csv(file_path, parse_dates=['time'])
46     except FileNotFoundError:
```

```

47     logging.error(f"File not found: {file_path}")
48     sys.exit(1)
49 except pd.errors.ParserError as e:
50     logging.error(f"Error parsing CSV file: {e}")
51     sys.exit(1)
52 except Exception as e:
53     logging.error(f"Unexpected error: {e}")
54     sys.exit(1)
55
56 rename_mapping = {
57     'time': 'Date',
58     'open': 'Open',
59     'high': 'High',
60     'low': 'Low',
61     'close': 'Close'
62 }
63 df.rename(columns=rename_mapping, inplace=True)
64
65 logging.info(f"Data columns after renaming: {df.columns.tolist()}")
66
67 df.sort_values('Date', inplace=True)
68 df.reset_index(drop=True, inplace=True)
69 logging.info("Data loaded and sorted successfully.")
70 return df
71
72
73 def compute_rsi(series, window=14):
74     delta = series.diff()
75     gain = delta.where(delta > 0, 0).rolling(window=window).mean()
76     loss = -delta.where(delta < 0, 0).rolling(window=window).mean()
77     RS = gain / (loss + 1e-9)
78     return 100 - (100 / (1 + RS))
79
80
81 def compute_macd(series, span_short=12, span_long=26, span_signal=9):
82     ema_short = series.ewm(span=span_short, adjust=False).mean()
83     ema_long = series.ewm(span=span_long, adjust=False).mean()
84     macd_line = ema_short - ema_long
85     signal_line = macd_line.ewm(span=span_signal, adjust=False).mean()
86     return macd_line - signal_line # histogram
87
88
89 def compute_obv(df):
90     signed_volume = (np.sign(df['Close'].diff()) * df['Volume']).fillna(0)
91     return signed_volume.cumsum()
92
93
94 def compute_adx(df, window=14):
95     """Pseudo-ADX approach using rolling True Range / Close."""
96     df['H-L'] = df['High'] - df['Low']
97     df['H-Cp'] = (df['High'] - df['Close'].shift(1)).abs()
98     df['L-Cp'] = (df['Low'] - df['Close'].shift(1)).abs()
99     tr = df[['H-L', 'H-Cp', 'L-Cp']].max(axis=1)
100     tr_rolling = tr.rolling(window=window).mean()
101
102     adx_placeholder = tr_rolling / (df['Close'] + 1e-9)
103     df.drop(['H-L', 'H-Cp', 'L-Cp'], axis=1, inplace=True)
104     return adx_placeholder
105
106
107 def compute_bollinger_bands(series, window=20, num_std=2):
108     sma = series.rolling(window=window).mean()
109     std = series.rolling(window=window).std()
110     upper = sma + num_std * std
111     lower = sma - num_std * std
112     bandwidth = (upper - lower) / (sma + 1e-9)
113     return upper, lower, bandwidth
114
115
116 def compute_mfi(df, window=14):
117     typical_price = (df['High'] + df['Low'] + df['Close']) / 3
118     money_flow = typical_price * df['Volume']
119     prev_tp = typical_price.shift(1)

```

```

120     flow_pos = money_flow.where(typical_price > prev_tp, 0)
121     flow_neg = money_flow.where(typical_price < prev_tp, 0)
122     pos_sum = flow_pos.rolling(window=window).sum()
123     neg_sum = flow_neg.rolling(window=window).sum()
124     mfi = 100 - (100 / (1 + pos_sum/(neg_sum+1e-9)))
125     return mfi
126
127
128 def calculate_technical_indicators(df):
129     logging.info("Calculating technical indicators...")
130
131     df['RSI'] = compute_rsi(df['Close'], window=14)
132     df['MACD'] = compute_macd(df['Close'])
133     df['OBV'] = compute_obv(df)
134     df['ADX'] = compute_adx(df)
135
136     upper_bb, lower_bb, bb_width = compute_bollinger_bands(df['Close'], window=20, num_std=2)
137     df['BB_Upper'] = upper_bb
138     df['BB_Lower'] = lower_bb
139     df['BB_Width'] = bb_width
140
141     df['MFI'] = compute_mfi(df, window=14)
142
143     df['SMA_5'] = df['Close'].rolling(window=5).mean()
144     df['SMA_10'] = df['Close'].rolling(window=10).mean()
145     df['EMA_5'] = df['Close'].ewm(span=5, adjust=False).mean()
146     df['EMA_10'] = df['Close'].ewm(span=10, adjust=False).mean()
147
148     df['STDDEV_5'] = df['Close'].rolling(window=5).std()
149     df.dropna(inplace=True)
150     logging.info("Technical indicators calculated successfully.")
151     return df
152
153
154 #####
155 # 2. ARGUMENT PARSING
156 #####
157 def parse_arguments():
158     parser = argparse.ArgumentParser(description='Train LSTM and DQN models for stock trading.')
159     parser.add_argument('csv_path', type=str, help='Path to the CSV data file (with columns time,open,high,low,close,volume).')
160     return parser.parse_args()
161
162
163 #####
164 # 3. MAIN
165 #####
166 def main():
167     # 1) Parse args
168     args = parse_arguments()
169     csv_path = args.csv_path
170
171     # 2) Load data & advanced indicators
172     data = load_data(csv_path)
173     data = calculate_technical_indicators(data)
174
175     # EXCLUDE 'Close' from feature inputs
176     feature_columns = [
177         'SMA_5', 'SMA_10', 'EMA_5', 'EMA_10', 'STDDEV_5',
178         'RSI', 'MACD', 'ADX', 'OBV', 'Volume', 'Open', 'High', 'Low',
179         'BB_Upper', 'BB_Lower', 'BB_Width', 'MFI'
180     ]
181     target_column = 'Close'
182     data = data[['Date'] + feature_columns + [target_column]]
183     data.dropna(inplace=True)
184
185     # 3) Scale
186     scaler_features = MinMaxScaler()
187     scaler_target = MinMaxScaler()
188
189     X_all = data[feature_columns].values
190     y_all = data[[target_column]].values

```

```

191 X_scaled = scaler_features.fit_transform(X_all)
192 y_scaled = scaler_target.fit_transform(y_all).flatten()
193
194
195 # 4) Create LSTM Sequences
196 def create_sequences(features, target, window_size=15):
197     X_seq, y_seq = [], []
198     for i in range(len(features) - window_size):
199         X_seq.append(features[i:i+window_size])
200         y_seq.append(target[i+window_size])
201     return np.array(X_seq), np.array(y_seq)
202
203 window_size = 15
204 X, y = create_sequences(X_scaled, y_scaled, window_size)
205
206 # 5) Train/Val/Test Split
207 train_size = int(len(X)*0.7)
208 val_size = int(len(X)*0.15)
209 test_size = len(X) - train_size - val_size
210
211 X_train = X[:train_size]
212 y_train = y[:train_size]
213 X_val = X[train_size:train_size+val_size]
214 y_val = y[train_size:train_size+val_size]
215 X_test = X[train_size+val_size:]
216 y_test = y[train_size+val_size:]
217
218 logging.info(f"Scaled training features shape: {X_train.shape}")
219 logging.info(f"Scaled validation features shape: {X_val.shape}")
220 logging.info(f"Scaled testing features shape: {X_test.shape}")
221 logging.info(f"Scaled training target shape: {y_train.shape}")
222 logging.info(f"Scaled validation target shape: {y_val.shape}")
223 logging.info(f"Scaled testing target shape: {y_test.shape}")
224
225 # 6) GPU/CPU Config
226 def configure_device():
227     gpus = tf.config.list_physical_devices('GPU')
228     if gpus:
229         try:
230             for gpu in gpus:
231                 tf.config.experimental.set_memory_growth(gpu, True)
232                 logging.info(f"{len(gpus)} GPU(s) detected and configured.")
233         except RuntimeError as e:
234             logging.error(e)
235     else:
236         logging.info("No GPU detected, using CPU.")
237
238 configure_device()
239
240 # 7) Build LSTM
241 def build_advanced_lstm(input_shape, hyperparams):
242     model = Sequential()
243     for i in range(hyperparams['num_lstm_layers']):
244         return_seqs = (i < hyperparams['num_lstm_layers'] - 1)
245         model.add(Bidirectional(
246             LSTM(hyperparams['lstm_units'],
247                 return_sequences=return_seqs,
248                 kernel_regularizer=tf.keras.regularizers.l2(0.001)
249             ), input_shape=input_shape if i==0 else None))
250         model.add(Dropout(hyperparams['dropout_rate']))
251
252     model.add(Dense(1, activation='linear'))
253
254 # Optimizer
255 if hyperparams['optimizer'] == 'Adam':
256     opt = Adam(learning_rate=hyperparams['learning_rate'], decay=hyperparams['decay'])
257 elif hyperparams['optimizer'] == 'Nadam':
258     opt = Nadam(learning_rate=hyperparams['learning_rate'])
259 else:
260     opt = Adam(learning_rate=hyperparams['learning_rate'])
261
262 model.compile(optimizer=opt, loss=Huber(), metrics=['mae'])
263 return model

```

```

264
265 # 8) Optuna Tuning
266 def objective(trial):
267     num_lstm_layers = trial.suggest_int('num_lstm_layers', 1, 3)
268     lstm_units = trial.suggest_categorical('lstm_units', [32, 64, 96, 128])
269     dropout_rate = trial.suggest_float('dropout_rate', 0.1, 0.5)
270     learning_rate = trial.suggest_loguniform('learning_rate', 1e-5, 1e-2)
271     optimizer_name = trial.suggest_categorical('optimizer', ['Adam', 'Nadam'])
272     decay = trial.suggest_float('decay', 0.0, 1e-4)
273
274     hyperparams = {
275         'num_lstm_layers': num_lstm_layers,
276         'lstm_units': lstm_units,
277         'dropout_rate': dropout_rate,
278         'learning_rate': learning_rate,
279         'optimizer': optimizer_name,
280         'decay': decay
281     }
282
283     model_ = build_advanced_lstm((X_train.shape[1], X_train.shape[2]), hyperparams)
284
285     early_stop = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
286     lr_reduce = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=5, min_lr=1e
287                                     -6)
288
289     cb_prune = KerasPruningCallback(trial, 'val_loss')
290
291     history = model_.fit(
292         X_train, y_train,
293         epochs=100,
294         batch_size=16,
295         validation_data=(X_val, y_val),
296         callbacks=[early_stop, lr_reduce, cb_prune],
297         verbose=0
298     )
299     val_mae = min(history.history['val_mae'])
300     return val_mae
301
302 logging.info("Starting hyperparameter optimization with Optuna...")
303 study = optuna.create_study(direction='minimize')
304 study.optimize(objective, n_trials=50) # might take a long time
305
306 best_params = study.best_params
307 logging.info(f"Best Hyperparameters from Optuna: {best_params}")
308
309 # 9) Train Best LSTM
310 best_model = build_advanced_lstm((X_train.shape[1], X_train.shape[2]), best_params)
311 early_stop2 = EarlyStopping(monitor='val_loss', patience=20, restore_best_weights=True)
312 lr_reduce2 = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=5, min_lr=1e-6)
313
314 logging.info("Training the best LSTM model with optimized hyperparameters...")
315 history = best_model.fit(
316     X_train, y_train,
317     epochs=300,
318     batch_size=16,
319     validation_data=(X_val, y_val),
320     callbacks=[early_stop2, lr_reduce2],
321     verbose=1
322 )
323
324 # 10) Evaluate
325 def evaluate_model(model, X_test, y_test):
326     logging.info("Evaluating model...")
327     # Predict scaled
328     y_pred_scaled = model.predict(X_test).flatten()
329     y_pred_scaled = np.clip(y_pred_scaled, 0, 1) # clamp if needed
330     # Inverse
331     y_pred = scaler_target.inverse_transform(y_pred_scaled.reshape(-1,1)).flatten()
332     y_test_actual = scaler_target.inverse_transform(y_test.reshape(-1,1)).flatten()
333
334     mse = mean_squared_error(y_test_actual, y_pred)
335     rmse = np.sqrt(mse)
336     mae = mean_absolute_error(y_test_actual, y_pred)

```

```

336     r2 = r2_score(y_test_actual, y_pred)
337
338     # Directional accuracy
339     direction_actual = np.sign(np.diff(y_test_actual))
340     direction_pred = np.sign(np.diff(y_pred))
341     directional_accuracy = np.mean(direction_actual == direction_pred)
342
343     logging.info(f"Test MSE: {mse}")
344     logging.info(f"Test RMSE: {rmse}")
345     logging.info(f"Test MAE: {mae}")
346     logging.info(f"Test R2 Score: {r2}")
347     logging.info(f"Directional Accuracy: {directional_accuracy}")
348
349     # Plot
350     plt.figure(figsize=(14, 7))
351     plt.plot(y_test_actual, label='Actual Price')
352     plt.plot(y_pred, label='Predicted Price')
353     plt.title('Actual vs Predicted Prices')
354     plt.xlabel('Time Step')
355     plt.ylabel('Price')
356     plt.legend()
357     plt.grid(True)
358     plt.savefig('actual_vs_predicted.png')
359     plt.close()
360     logging.info("Actual vs Predicted plot saved as 'actual_vs_predicted.png'")
361
362     # Tabulate first 40 predictions
363     table_data = []
364     for i in range(min(40, len(y_test_actual))):
365         table_data.append([i, round(y_test_actual[i],2), round(y_pred[i],2)])
366     headers = ["Index", "Actual Price", "Predicted Price"]
367     print(tabulate(table_data, headers=headers, tablefmt="pretty"))
368
369     return mse, rmse, mae, r2, directional_accuracy
370
371 mse, rmse, mae, r2, directional_accuracy = evaluate_model(best_model, X_test, y_test)
372
373 # 11) Save
374 best_model.save('optimized_lstm_model.h5')
375 import joblib
376 joblib.dump(scaler_features, 'scaler_features.save')
377 joblib.dump(scaler_target, 'scaler_target.save')
378 logging.info("Model and scalers saved as 'optimized_lstm_model.h5', 'scaler_features.save'
379             ', and 'scaler_target.save'.")
380
381 #####
382 # 12) Reinforcement Learning Environment
383 #####
384 class StockTradingEnv(gym.Env):
385     """
386     A simple stock trading environment for OpenAI Gym
387     """
388     metadata = {'render.modes': ['human']}
389
390     def __init__(self, df, initial_balance=10000):
391         super().__init__()
392         self.df = df.reset_index()
393         self.initial_balance = initial_balance
394         self.balance = initial_balance
395         self.net_worth = initial_balance
396         self.max_steps = len(df)
397         self.current_step = 0
398         self.shares_held = 0
399         self.cost_basis = 0
400
401         # We re-use feature_columns from above
402         # (Excluding 'Close' from the observation)
403         # Actions: 0=Sell, 1=Hold, 2=Buy
404         self.action_space = spaces.Discrete(3)
405
406         # Observations => advanced feature columns + 3 additional (balance, shares,
407         cost_basis)
408         self.observation_space = spaces.Box(

```

```

407         low=0,
408         high=1,
409         shape=(len(feature_columns) + 3,),
410         dtype=np.float32
411     )
412
413     def reset(self):
414         self.balance = self.initial_balance
415         self.net_worth = self.initial_balance
416         self.current_step = 0
417         self.shares_held = 0
418         self.cost_basis = 0
419         return self._next_observation()
420
421     def _next_observation(self):
422         obs = self.df.loc[self.current_step, feature_columns].values
423         # Simple normalization by max to keep it [0,1]
424         obs = obs / np.max(obs) if np.max(obs) != 0 else obs
425
426         additional = np.array([
427             self.balance / self.initial_balance,
428             self.shares_held / 100.0,
429             self.cost_basis / self.initial_balance
430         ])
431         return np.concatenate([obs, additional])
432
433     def step(self, action):
434         current_price = self.df.loc[self.current_step, 'Close']
435
436         if action == 2: # Buy
437             total_possible = self.balance // current_price
438             shares_bought = total_possible
439             if shares_bought > 0:
440                 self.balance -= shares_bought * current_price
441                 self.shares_held += shares_bought
442                 self.cost_basis = (
443                     (self.cost_basis * (self.shares_held - shares_bought)) +
444                     (shares_bought * current_price)
445                 ) / self.shares_held
446
447         elif action == 0: # Sell
448             if self.shares_held > 0:
449                 self.balance += self.shares_held * current_price
450                 self.shares_held = 0
451                 self.cost_basis = 0
452
453         self.net_worth = self.balance + self.shares_held * current_price
454         self.current_step += 1
455
456         done = (self.current_step >= self.max_steps - 1)
457         reward = self.net_worth - self.initial_balance
458
459         obs = self._next_observation()
460         return obs, reward, done, {}
461
462     def render(self, mode='human'):
463         profit = self.net_worth - self.initial_balance
464         print(f"Step: {self.current_step}")
465         print(f"Balance: {self.balance}")
466         print(f"Shares held: {self.shares_held} (Cost Basis: {self.cost_basis})")
467         print(f"Net worth: {self.net_worth}")
468         print(f"Profit: {profit}")
469
470     def train_dqn_agent(env):
471         logging.info("Training DQN Agent...")
472         try:
473             model = DQN(
474                 'MlpPolicy',
475                 env,
476                 verbose=1,
477                 learning_rate=1e-3,
478                 buffer_size=10000,
479                 learning_starts=1000,

```



```

480         batch_size=64,
481         tau=1.0,
482         gamma=0.99,
483         train_freq=4,
484         target_update_interval=1000,
485         exploration_fraction=0.1,
486         exploration_final_eps=0.02,
487         tensorboard_log="./dqn_stock_tensorboard/"
488     )
489     model.learn(total_timesteps=100000)
490     model.save("dqn_stock_trading")
491     logging.info("DQN Agent trained and saved as 'dqn_stock_trading.zip'.")
492     return model
493 except Exception as e:
494     logging.error(f"Error training DQN Agent: {e}")
495     sys.exit(1)
496
497 # Initialize RL environment
498 logging.info("Initializing and training DQN environment...")
499 trading_env = StockTradingEnv(data)
500 trading_env = DummyVecEnv([lambda: trading_env])
501
502 # Train
503 dqn_model = train_dqn_agent(trading_env)
504
505 logging.info("All tasks complete. Exiting.")
506
507
508 if __name__ == "__main__":
509     main()

```

3 Step-by-Step Explanation

3.1 Imports and Configuration

- numpy, pandas, sklearn are used for data manipulation and scaling.
- matplotlib and seaborn are used for plotting.
- tensorflow.keras is used for building the LSTM model.
- optuna is used for hyperparameter tuning.
- gym and stable_baselines3 are used for the reinforcement learning (RL) portion of the code.
- The environment variable TF_CPP_MIN_LOG_LEVEL suppresses TensorFlow messages to keep logs clean.

3.2 Data Loading and Preprocessing (Functions in Step 1)

1. load_data(file_path):

- Reads a CSV file, expecting columns time, open, high, low, close, volume.
- Renames these columns to Date, Open, High, Low, Close.
- Sorts the DataFrame by date, resets the index, and returns the processed DataFrame.

2. Technical Indicators Functions: The code defines several functions to compute widely used technical indicators:

- compute_rsi(series, window=14): Calculates the Relative Strength Index over a given window.
- compute_macd(series, ...): Computes the Moving Average Convergence Divergence (MACD) histogram.
- compute_obv(df): Computes On-Balance Volume.
- compute_adx(df, window=14): Returns a placeholder ADX-like measure.

- `compute_bollinger_bands(series, ...)`: Returns Bollinger Bands and the band-width.
- `compute_mfi(df, window=14)`: Computes the Money Flow Index.

3. `calculate_technical_indicators(df)`:

- Applies all of the above functions on the `df`.
- Creates short/medium moving averages like SMA, EMA.
- Removes any resulting NaN values from the DataFrame.
- Returns the DataFrame with all computed indicators.

3.3 Argument Parsing (Step 2)

- Uses `argparse` to read a single argument, `csv_path`, which specifies the path to the CSV data file.

3.4 Main Function (Step 3)

The **main** function orchestrates the entire workflow:

1. **Parse arguments**: Calls `parse_arguments()` to get `csv_path`.
2. **Load and process data**:
 - Calls `load_data(csv_path)` to load the CSV.
 - Calls `calculate_technical_indicators(data)` to augment it with technical features.
3. **Setup feature and target columns**: The code excludes `Close` from the features (since it is the prediction target) and identifies the target as the `Close` column.
4. **Scaling**:
 - Instantiates two `MinMaxScalers`: one for features and one for the target.
 - Fits and transforms the data to a 0-1 range.
5. **Sequence Creation** (`create_sequences`):
 - Converts a time series of scaled features/targets into 3D sequences (`samples × window_size × number_of_features`).
 - Correspondingly, the targets become a vector of length (`original_length - window_size`).
6. **Train-Validation-Test Split**: Splits the resulting 3D array (and 1D targets) as 70% train, 15% val, and 15% test.
7. **GPU/CPU Configuration**: The function `configure_device()` attempts to set memory growth on any detected GPUs.

3.5 Building and Tuning the LSTM (Steps 7–9)

1. `build_advanced_lstm`:

- Takes `input_shape` (the shape of the input sequence) and a `hyperparams` dictionary.
- Builds a stacked or repeated LSTM architecture, each wrapped in a `Bidirectional` layer, with dropout in between.
- A final Dense layer of size 1 is added to output the price prediction.
- Compilation uses Huber loss and one of the possible optimizers (Adam or Nadam) with various learning rate/decay settings.

2. **Optuna Tuning** (`objective function`):

- Suggests a search space for the number of LSTM layers, units, dropout rate, learning rate, etc.
- Builds and trains the model with early stopping, learning rate reduction, and a pruning callback.
- Returns the best validation MAE found.
- This tuning is done in `study.optimize(..., n_trials=50)`.

3. Train the best model:

- Retrieves the best hyperparameters from the study, builds a new model, and trains it again for up to 300 epochs (with patience).

3.6 Evaluation and Saving the Model (Steps 10–11)

1. Evaluation:

- Predicts on test data.
- Inverse-transforms these predictions using `scaler_target`.
- Computes MSE, RMSE, MAE, R^2 , and a directional accuracy metric.
- Saves a plot of Actual vs. Predicted.
- Prints a table of the first 40 predictions for inspection.

2. Saving:

- Saves the trained model as `optimized_lstm_model.h5`.
- Saves the fitted `MinMaxScalers` via `joblib`.

3.7 Reinforcement Learning Environment and DQN Training (Step 12)

• `StockTradingEnv`:

- Inherits from `gym.Env`.
- Observations are a combination of scaled technical features and some additional info (balance ratio, shares held, cost basis).
- Action space has three discrete actions: Sell (0), Hold (1), and Buy (2).
- `step(action)` adjusts balance, shares, and net worth accordingly and returns a reward of `net_worth - initial_balance`.

• `train_dqn_agent(env)`:

- Creates a DQN model with MLP policy from `stable_baselines3`.
- Trains for 100,000 timesteps.
- Saves the resulting model to disk as `dqn_stock_trading`.

4 Replacing DQN with PPO

In many cases, **Proximal Policy Optimization (PPO)** can be more robust or sample-efficient than DQN for continuous or more complex action spaces (and sometimes even for discrete spaces). It uses a different approach to update the policy in a clipped fashion, preserving stability. Here is a step-by-step guide on how to replace the current DQN approach with PPO in your code:

4.1 Potential Benefits of PPO over DQN

1. **Policy Gradient Method:** PPO is an on-policy algorithm that directly optimizes a policy. This can be more stable in some environments.
2. **Clipping Mechanism:** PPO uses a clipping in the objective function that prevents large updates from one iteration to the next, making training more stable.
3. **Sample Efficiency:** In certain setups, PPO can reuse on-policy samples effectively, especially if you carefully manage the rollout steps.

4.2 Modifications in the Code

To make the switch, you only need to modify a few lines where DQN is used:

1. Import PPO:

```
1 # From:
2 from stable_baselines3 import DQN
3
4 # To:
5 from stable_baselines3 import PPO
```

2. Create the PPO Model Instead of DQN:

```
1 def train_ppo_agent(env):
2     logging.info("Training PPO Agent...")
3     try:
4         model = PPO(
5             'MlpPolicy',
6             env,
7             verbose=1,
8             learning_rate=3e-4, # typical default for PPO
9             n_steps=2048,      # number of steps to run for each environment
10            batch_size=64,
11            gamma=0.99,
12            gae_lambda=0.95,
13            clip_range=0.2,
14            ent_coef=0.0,
15            tensorboard_log="./ppo_stock_tensorboard/"
16        )
17        model.learn(total_timesteps=100000)
18        model.save("ppo_stock_trading")
19        logging.info("PPO Agent trained and saved as 'ppo_stock_trading.zip'.")
20        return model
21    except Exception as e:
22        logging.error(f"Error training PPO Agent: {e}")
23        sys.exit(1)
```

3. Instantiate the PPO Training:

- Instead of calling `train_dqn_agent`, call `train_ppo_agent`:

```
1 logging.info("Initializing and training PPO environment...")
2 trading_env = StockTradingEnv(data)
3 trading_env = DummyVecEnv([lambda: trading_env])
4
5 # Train
6 ppo_model = train_ppo_agent(trading_env)
```

Everything else regarding the environment remains the same. The main difference is in how the policy is updated and how the rollout buffer is handled internally.

5 Conclusion

By following the explanations in this document, you should now:

- Understand the structure and purpose of each code block in the LSTM + DQN trading workflow.
- Know how to easily switch from DQN to PPO by changing the import statements and the agent creation and training code.
- Appreciate some of the benefits (and potential pitfalls) of using PPO instead of DQN for reinforcement learning tasks in stock trading environments.

Disclaimer: Reinforcement learning models in algorithmic trading should be carefully tested under realistic constraints, including transaction costs, slippage, and risk management rules. Past performance does not guarantee future results.